

DATA SCIENTIST JOB RECOMMENDER SYSTEM

- Names:
 - Ingavi Kilavuka
 - Calvin Omwega
 - Alvin Kimathi
 - Ronny Kabiru
 - Carol Mundia
- Instructor: Maryann Mwikali
- Modeling Focus: Hybrid NLP & Recommender system

1. Business Understanding

Objective

In the pursuit for jobs in our desired field, This project aims to:

1. *Enhance Job Matching Precision:* Provide personalized job recommendations to users based on their input, focusing on job titles or descriptions that closely match the user's profile.
2. *Improve User Engagement:* By offering relevant job suggestions, increase user interaction with the platform, leading to higher conversion rates (e.g., more job applications).
3. *Expand Service Offerings:* Add more features to the recommendation system, such as recommendations based on salary expectations, location preferences, or industry-specific needs.
4. *Increase Platform Reach:* Position the job recommender system as a unique selling point to attract more users to the platform.

Problem Statement:

In the rapidly evolving job market, especially within the data science and analytics field, users are often overwhelmed by the vast number of job opportunities available. Many struggle to find positions that align with their specific skills, experience, and preferences in this specialized domain. The problem is exacerbated by traditional job platforms, which often lack the capability to provide personalized recommendations tailored to the unique demands of data science and analytics roles. This project aims to address this gap by developing a job recommender system that helps users find data science and analytics jobs that closely match their job title or description preferences, thereby increasing the chances of finding a suitable position more quickly.

Imports

```
import pandas as pd
import numpy as np
```

```

import re
import nltk
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.tokenize import word_tokenize
nltk.download('punkt')
nltk.download('stopwords')
from collections import Counter
from surprise import SVD, Dataset, Reader, accuracy
from surprise.model_selection import GridSearchCV, train_test_split,
cross_validate
from surprise import accuracy
import matplotlib.pyplot as plt
import seaborn as sns
! pip install surprise
from surprise import SVD, KNNBasic, KNNWithMeans, Reader,
Dataset, KNNBaseline
from surprise.model_selection import cross_validate, train_test_split
from surprise import accuracy
from surprise.model_selection import cross_validate
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from surprise.prediction_algorithms import SVD
from surprise.prediction_algorithms import KNNWithMeans, KNNBasic,
KNNBaseline
from surprise.model_selection import GridSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer
import joblib
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.metrics.pairwise import linear_kernel
from surprise import NMF
from surprise.model_selection import cross_validate
import matplotlib.pyplot as plt
np.random.seed(42)

```

```

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\HomePC\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\HomePC\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!

```

```

Requirement already satisfied: surprise in c:\users\homepc\anaconda3\
envs\learn-env\lib\site-packages (0.1)
Requirement already satisfied: scikit-surprise in c:\users\homepc\
anaconda3\envs\learn-env\lib\site-packages (from surprise) (1.1.1)
Requirement already satisfied: joblib>=0.11 in c:\users\homepc\
anaconda3\envs\learn-env\lib\site-packages (from scikit-surprise-

```

```
>surprise) (1.4.2)
Requirement already satisfied: numpy>=1.11.2 in c:\users\homepc\
anaconda3\envs\learn-env\lib\site-packages (from scikit-surprise-
>surprise) (1.22.0)
Requirement already satisfied: scipy>=1.0.0 in c:\users\homepc\
anaconda3\envs\learn-env\lib\site-packages (from scikit-surprise-
>surprise) (1.5.0)
Requirement already satisfied: six>=1.10.0 in c:\users\homepc\
anaconda3\envs\learn-env\lib\site-packages (from scikit-surprise-
>surprise) (1.15.0)
```

```
[notice] A new release of pip is available: 24.1.2 -> 24.2
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
df = pd.read_csv('jd_structured_data.csv')
df.columns
Index(['Job Title', 'Rating', 'Company Name', 'Location',
      'Headquarters',
      'Size', 'Founded', 'Type of ownership', 'Industry', 'Sector',
      'Competitors', 'Average Salary', 'Average Revenue',
      'Processed_JD'],
      dtype='object')
df.Location.unique()
array(['Albuquerque, NM', 'Linthicum, MD', 'Clearwater, FL',
      'Richland, WA', 'New York, NY', 'Dallas, TX', 'Baltimore, MD',
      'San Jose, CA', 'Rochester, NY', 'Chantilly, VA', 'Plano, TX',
      'Seattle, WA', 'Cambridge, MA', 'Newark, NJ', 'Mountain View,
CA',
      'San Francisco, CA', 'Denver, CO', 'Chicago, IL', 'Louisville,
KY',
      'Oregon', 'Herndon, VA', 'Hillsboro, OR', 'Worcester, MA',
      'Groton, CT', 'Detroit, MI', 'Sunnyvale, CA', 'Ipswich, MA',
      'Redlands, CA', 'Woburn, MA', 'Fremont, CA', 'Long Beach, NY',
      'Marlborough, MA', 'Allendale, NJ', 'Chandler, AZ',
      'Washington, DC', 'Bellevue, WA', 'Longmont, CO',
      'Beavercreek, OH', 'Peoria, IL', 'Kingdom, IL',
      'Fort Lauderdale, FL', 'Boston, MA', 'Huntsville, AL',
      'Armonk, NY', 'San Diego, CA', 'Saint Louis, MO', 'Lincoln,
RI',
      'Cincinnati, OH', 'Palo Alto, CA', 'Coraopolis, PA',
      'Framingham, MA', 'Atlanta, GA', 'New Jersey', 'Philadelphia,
PA',
      'Vancouver, WA', 'Indianapolis, IN', 'Lake Forest, IL',
      'Maryland Heights, MO', 'Charlottesville, VA', 'Pittsburgh,
```

PA',
'Harrisburg, PA', 'Laurel, MD', 'Arlington, VA', 'Tacoma, WA',
'Miami, FL', 'New Orleans, LA', 'Landover, MD',
'Patuxent River, MD', 'Suitland, MD', 'McLean, VA',
'Fort Belvoir, VA', 'Milwaukee, WI', 'Silver Spring, MD',
'South San Francisco, CA', 'Syracuse, NY', 'Gainesville, FL',
'Virginia', 'Houston, TX', 'Charlotte, NC', 'Southfield, MI',
'Matawan, NJ', 'Rochester, MN', 'Phoenix, AZ', 'Omaha, NE',
'Lyndhurst, NJ', 'Atlanta, IN', 'Rockville, MD', 'Minneapolis',

MN',
'Santa Barbara, CA', 'Bethesda, MD', 'Los Angeles, CA',
'Alabaster, AL', 'Mount Olive, NJ',
'Santa Fe Springs, Los Angeles, CA', 'Kansas City, MO',
'Ashburn, VA', 'Fort Worth, TX', 'Valencia, CA', 'Novato, CA',
'Reston, VA', 'Aurora, CO', 'Waltham, MA', 'Tampa, FL',
'Camarillo, CA', 'Riverton, UT', 'Chattanooga, TN', 'Milpitas',

CA',
'Ewing, NJ', 'Cupertino, CA', 'Alexandria, VA', 'Frederick',

MD',
'Fort Lee, NJ', 'West Reading, PA', 'Madison, WI', 'Dearborn',

MI',
'Winter Park, FL', 'San Rafael, CA', 'Hamilton, NJ',
'Woodbridge, NJ', 'Lewes, DE', 'Springfield, MO', 'Burbank',

CA',
'Newton, MA', 'Salt Lake City, UT', 'Lafayette, LA',
'Annapolis Junction, MD', 'Carlsbad, CA', 'Highland, CA',
'Burleson, TX', 'Hoopeston, IL', 'Georgia Southern, GA',
'Scotts Valley, CA', 'Knoxville, TN', 'Millville, DE',
'Sheboygan, WI', 'San Mateo, CA', 'Dayton, OH', 'Parlier, CA',
'Meridian, ID', 'Cherry Hill, NJ', 'Columbia, MD', 'Anchorage',

AK',
'Lincoln, NE', 'Nashville, TN', 'Portland, OR', 'Hoboken, NJ',
'Port Washington, NY', 'Austin, TX', 'Grand Rapids, MI',
'Creve Coeur, MO', 'United States', 'Providence, RI',
'Raleigh, NC', 'Phila, PA', 'Oakland, CA', 'Bloomington, MN',
'Boise, ID', 'Oak Ridge, TN', 'Agoura Hills, CA', 'Pella, IA',
'Burlington, MA', 'Woodinville, WA', 'San Ramon, CA',
'Red Bank, NJ', 'Pomona, CA', 'Columbia, SC', 'Springfield',

MA',
'San Antonio, TX', 'Portsmouth, VA', 'West Palm Beach, FL',
'Newark, CA', 'Exton, PA', 'Owensboro, KY', 'Hartford, CT',
'Orange, CA', 'Tallahassee, FL', 'Lenexa, KS', 'Concord, CA',
'Vail, CO', 'Dulles, VA', 'Natick, MA', 'Winston-Salem, NC',
'Richfield, OH', 'Hampton, VA', 'Ithaca, NY', 'Marietta, GA',
'Quincy, MA', 'Green Bay, WI', 'Durham, NC', 'Clovis, CA',
'Orlando, FL', 'Columbia, MO', 'Fountain Valley, CA',
'Westlake, OH', 'Des Moines, IA', 'Bedford, MA',
'Cedar Rapids, IA', 'Blue Bell, PA', 'Springfield, VA',
'Jersey City, NJ', 'Emeryville, CA', 'Carle Place, NY',

```
'King of Prussia, PA', 'Santa Clara, CA', 'Piscataway, NJ',  
'Brisbane, CA', 'Foster City, CA', 'Holyoke, MA', 'Juncos, PR',  
'Corvallis, OR', 'Gaithersburg, MD', 'Aliso Viejo, CA',  
'Dublin, CA', 'Bristol, TN', 'Arvada, CO', 'Franklin, TN',  
'Plymouth Meeting, PA', 'Allentown, PA', 'Logan, UT',  
'Birmingham, AL', 'Scottsdale, AZ', 'Bloomington, IL',  
'Alameda, CA', 'Johns Creek, GA', 'Roanoke, VA', 'Maryland',  
'Glen Burnie, MD', 'Watertown, MA', 'Cambridge, MD',  
'Amsterdam, NY', 'Irvine, CA', 'Ann Arbor, MI', 'Sandy, UT',  
'Olympia, WA', 'Richmond, VA', 'Tempe, AZ', 'Michigan']  
dtype=object)
```

```
df.shape
```

```
(956, 14)
```

```
df.describe()
```

	Rating	Size	Founded	Average Salary	Average
Revenue					
count	956.000000	956.000000	956.000000	956.000000	956.000000
mean	3.601255	3027.393199	1774.605649	103.153900	24319.000761
std	1.067619	3677.688565	598.942517	31.971932	60571.308570
min	-1.000000	-1.000000	-1.000000	15.500000	1.000000
25%	3.300000	350.500000	1937.000000	84.500000	17.500000
50%	3.800000	750.500000	1992.000000	103.153900	1500.000000
75%	4.200000	3027.393199	2008.000000	114.000000	24319.000761
max	5.000000	10000.000000	2019.000000	254.000000	250500.000000

```
df.dtypes
```

Job Title	object
Rating	float64
Company Name	object
Location	object
Headquarters	object
Size	float64
Founded	int64
Type of ownership	object
Industry	object
Sector	object
Competitors	object
Average Salary	float64

Average Revenue float64
Processed_JD object
dtype: object

df

		Job Title	Rating \
0		Data Scientist	3.8
1	Healthcare	Data Scientist	3.4
2		Data Scientist	4.8
3		Data Scientist	3.8
4		Data Scientist	2.9
...	
951		Senior Data Engineer	4.4
952	Project Scientist - Auton Lab, Robotics Institute		2.6
953		Data Science Manager	3.2
954		Data Engineer	4.8
955	Research Scientist – Security and Privacy		3.6

	Company Name	Location \
0	Tecolote Research	Albuquerque, NM
1	University of Maryland Medical System	Linthicum, MD
2	KnowBe4	Clearwater, FL
3	PNNL	Richland, WA
4	Affinity Solutions	New York, NY
...
951	Eventbrite	Nashville, TN
952	Software Engineering Institute	Pittsburgh, PA
953	Numeric, LLC	Allentown, PA
954	IGNW	Austin, TX
955	Riverside Research Institute	Beavercreek, OH

	Headquarters	Size	Founded	Type of ownership \
0	Goleta, CA	750.5	1973	Company - Private
1	Baltimore, MD	10000.0	1984	Other Organization
2	Clearwater, FL	750.5	2010	Company - Private
3	Richland, WA	3000.5	1965	Government
4	New York, NY	125.5	1998	Company - Private
...
951	San Francisco, CA	3000.5	2006	Company - Public
952	Pittsburgh, PA	750.5	1984	College / University
953	Chadds Ford, PA	25.5	-1	Company - Private
954	Portland, OR	350.5	2015	Company - Private
955	Arlington, VA	750.5	1967	Nonprofit Organization

	Industry	Sector \
0	Aerospace & Defense	Aerospace & Defense
1	Health Care Services & Hospitals	Health Care
2	Security Services	Business Services
3	Energy	Oil, Gas, Energy & Utilities

4	Advertising & Marketing	Business Services
..
951	Internet	Information Technology
952	Colleges & Universities	Education
953	Staffing & Outsourcing	Business Services
954	IT Services	Information Technology
955	Federal Agencies	Government
		Competitors Average Salary
\		
0		-1 72.0000
1		-1 87.5000
2		-1 85.0000
3	Oak Ridge National Laboratory, National Renewa...	76.5000
4	Commerce Signals, Cardlytics, Yodlee	114.5000
..
951	See Tickets, TicketWeb, Vendini	102.5000
952		-1 73.5000
953		-1 127.5000
954	Slalom	103.1539
955		-1 93.5000
Average Revenue		
Processed_JD		
0	75.000000	Data Scientist Location: Albuquerque, Educatio...
1	3500.000000	What You Will Do: General Summary The Healthca...
2	300.000000	KnowBe4, Inc. high growth information security...
3	250500.000000	*Organization Job ID** Job ID: 310709 Director...
4	24319.000761	Data Scientist Affinity Solutions Marketing Cl...
..
.		
951	300.000000	THE CHALLENGE Eventbrite world-class data repo...
952	24319.000761	The Auton Lab Carnegie Mellon University

```

large...
953          7.500000  Data Science ManagerResponsibilities:
Oversee ...
954          37.500000  Loading... Title: Data Engineer Location:
Aust...
955          75.000000  Returning Candidate? Log back Career Portal
cl...

```

```
[956 rows x 14 columns]
```

```

missing_percentage = df.isnull().mean() * 100
print(missing_percentage)

```

```

Job Title          0.0
Rating             0.0
Company Name       0.0
Location           0.0
Headquarters       0.0
Size               0.0
Founded            0.0
Type of ownership  0.0
Industry           0.0
Sector             0.0
Competitors        0.0
Average Salary     0.0
Average Revenue    0.0
Processed_JD       0.0
dtype: float64

```

Data Cleaning

```

df = df.drop('Competitors', axis=1)

missing_percentage = df.isnull().mean() * 100
print(missing_percentage)

```

```

Job Title          0.0
Rating             0.0
Company Name       0.0
Location           0.0
Headquarters       0.0
Size               0.0
Founded            0.0
Type of ownership  0.0
Industry           0.0
Sector             0.0
Average Salary     0.0
Average Revenue    0.0
Processed_JD       0.0
dtype: float64

```



```

df['Size'] = pd.to_numeric(df['Size'], errors='coerce')
df['Average Salary'] = pd.to_numeric(df['Average Salary'],
errors='coerce')
df['Average Revenue'] = pd.to_numeric(df['Average Revenue'],
errors='coerce')

q1 = df['Average Salary'].quantile(0.25)
q3 = df['Average Salary'].quantile(0.75)
iqr = q3 - q1

lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr
df = df[(df['Average Salary'] >= lower_bound) & (df['Average Salary']
<= upper_bound)]

text_columns = ['Job Title', 'Company Name', 'Location',
'Headquarters', 'Type of ownership', 'Industry', 'Sector',
'Processed_JD']
for col in text_columns:
    df[col] = df[col].str.strip().str.lower()

duplicates = df[df.duplicated()]
print(duplicates)

```

	Job Title	Rating	\
30	data scientist	4.8	
31	data scientist	3.8	
62	data scientist	4.1	
63	data scientist	3.4	
94	staff data scientist - technology	3.2	
..	
951	senior data engineer	4.4	
952	project scientist - auton lab, robotics institute	2.6	
953	data science manager	3.2	
954	data engineer	4.8	
955	research scientist - security and privacy	3.6	

	Company Name	Location	
Headquarters \			
30	knowbe4	clearwater, fl	clearwater,
fl			
31	pnnl	richland, wa	richland,
wa			
62	clearone advantage	baltimore, md	baltimore,
md			
63	cyrusone	dallas, tx	dallas,
tx			
94	walmart	plano, tx	bentonville,
ar			
..

951	eventbrite	nashville, tn	san francisco, ca
952	software engineering institute	pittsburgh, pa	pittsburgh, pa
953	numeric, llc	allentown, pa	chadds ford, pa
954	ignw	austin, tx	portland, or
955	riverside research institute	beavercreek, oh	arlington, va

	Size	Founded	Type of ownership \
30	750.5	2010	company - private
31	3000.5	1965	government
62	750.5	2008	company - private
63	350.5	2000	company - public
94	10000.0	1962	company - public
..
951	3000.5	2006	company - public
952	750.5	1984	college / university
953	25.5	-1	company - private
954	350.5	2015	company - private
955	750.5	1967	nonprofit organization

	Industry	Sector
30	security services	business services
31	energy	oil, gas, energy & utilities
62	banks & credit unions	finance
63	real estate	real estate
94	department, clothing, & shoe stores	retail
..
951	internet	information technology
952	colleges & universities	education
953	staffing & outsourcing	business services
954	it services	information technology
955	federal agencies	government

Average Salary	Average Revenue \
----------------	-------------------

30	85.0000	300.000000
31	76.5000	250500.000000
62	73.5000	24319.000761
63	95.0000	1500.000000
94	139.0000	10.000000
..
951	102.5000	300.000000
952	73.5000	24319.000761
953	127.5000	7.500000
954	103.1539	37.500000
955	93.5000	75.000000

```

Processed_JD
30 knowbe4, inc. high growth information security...
31 *organization job id** job id: 310709 director...
62 job description **please local candidates appl...
63 cyrusone seeking talented data scientist holds...
94 position summary... drives execution multiple ...
..
951 the challenge eventbrite world-class data repo...
952 the auton lab carnegie mellon university large...
953 data science managerresponsibilities: oversee ...
954 loading... title: data engineer location: aust...
955 returning candidate? log back career portal cl...

```

[334 rows x 13 columns]

```

df['Location'] = df['Location'].str.replace('new york, ny', 'new york,
new york')
df

```

	Job Title	Rating	\
0	data scientist	3.8	
1	healthcare data scientist	3.4	
2	data scientist	4.8	
3	data scientist	3.8	
4	data scientist	2.9	
..	
951	senior data engineer	4.4	
952	project scientist - auton lab, robotics institute	2.6	
953	data science manager	3.2	
954	data engineer	4.8	
955	research scientist – security and privacy	3.6	

	Company Name	Location	\
0	tecolote research	albuquerque, nm	
1	university of maryland medical system	linthicum, md	
2	knowbe4	clearwater, fl	
3	pnnl	richland, wa	
4	affinity solutions	new york, new york	

951	eventbrite	nashville, tn
952	software engineering institute	pittsburgh, pa
953	numeric, llc	allentown, pa
954	ignw	austin, tx
955	riverside research institute	beavercreek, oh

	Headquarters	Size	Founded	Type of ownership \
0	goleta, ca	750.5	1973	company - private
1	baltimore, md	10000.0	1984	other organization
2	clearwater, fl	750.5	2010	company - private
3	richland, wa	3000.5	1965	government
4	new york, ny	125.5	1998	company - private
951	san francisco, ca	3000.5	2006	company - public
952	pittsburgh, pa	750.5	1984	college / university
953	chadds ford, pa	25.5	-1	company - private
954	portland, or	350.5	2015	company - private
955	arlington, va	750.5	1967	nonprofit organization

	Industry	Sector \
0	aerospace & defense	aerospace & defense
1	health care services & hospitals	health care
2	security services	business services
3	energy	oil, gas, energy & utilities
4	advertising & marketing	business services
951	internet	information technology
952	colleges & universities	education
953	staffing & outsourcing	business services
954	it services	information technology
955	federal agencies	government

	Average Salary	Average Revenue \
0	72.0000	75.000000
1	87.5000	3500.000000
2	85.0000	300.000000
3	76.5000	250500.000000
4	114.5000	24319.000761
951	102.5000	300.000000
952	73.5000	24319.000761
953	127.5000	7.500000
954	103.1539	37.500000
955	93.5000	75.000000

	Processed_JD
0	data scientist location: albuquerque, educatio...
1	what you will do: general summary the healthca...
2	knowbe4, inc. high growth information security...

```

3    *organization job id** job id: 310709 director...
4    data scientist affinity solutions marketing cl...
..
951 the challenge eventbrite world-class data repo...
952 the auton lab carnegie mellon university large...
953 data science managerresponsibilities: oversee ...
954 loading... title: data engineer location: aust...
955 returning candidate? log back career portal cl...

[895 rows x 13 columns]

```

Data Engineering(NLP)

```

stop_words = set(stopwords.words('english'))

def remove_stopwords(text):
    if isinstance(text, str): # Check if the input is a string
        return ' '.join([word for word in text.split() if word.lower()
not in stop_words])
    else:
        return text # If it's not a string, return it as is (could be
None or some other type)

# Apply the function to the cleaned_text column
df['Processed_JD'] = df['Processed_JD'].apply(remove_stopwords)
df['Location'] = df['Location'].apply(remove_stopwords)
df['Sector'] = df['Sector'].apply(remove_stopwords)
df['Industry'] = df['Industry'].apply(remove_stopwords)
df['Location'] = df['Location'].apply(remove_stopwords)

# Display the first few rows to verify the changes
df[['Processed_JD', 'Location', 'Industry', 'Sector']].head()

```

	Location \	Processed_JD
0	data scientist location: albuquerque, educatio... nm	albuquerque,
1	do: general summary healthcare data scientist ... md	linthicum,
2	knowbe4, inc. high growth information security... fl	clearwater,
3	*organization job id** job id: 310709 director... wa	richland,
4	data scientist affinity solutions marketing cl... york	new york, new

	Industry	Sector
0	aerospace & defense	aerospace & defense
1	health care services & hospitals	health care
2	security services	business services

```

3         energy oil, gas, energy & utilities
4         advertising & marketing          business services

def tokenize_text(text):
    if isinstance(text, str): # Check if the input is a string
        return word_tokenize(text)
    else:
        return [] # Return an empty list if the text is None or not a
string

# Apply the function to the cleaned_text column
df['JD_tokens'] = df['Processed_JD'].apply(tokenize_text)
df['Location_tokens'] = df['Location'].apply(tokenize_text)
df['Industry_tokens'] = df['Industry'].apply(tokenize_text)
df['Sector_tokens'] = df['Sector'].apply(tokenize_text)

# Display the first few rows to verify the changes
df[['JD_tokens', 'Industry_tokens', 'Sector_tokens', 'Location_tokens']].head()

```


	JD_tokens \
0	[data, scientist, location, :, albuquerque, ,,,]
1	[do, :, general, summary, healthcare, data, sc...]
2	[knowbe4, ,, inc., high, growth, information, ...]
3	[* , organization, job, id, *, *, job, id, :, 3...]
4	[data, scientist, affinity, solutions, marketi...]

	Industry_tokens \
0	[aerospace, &, defense]
1	[health, care, services, &, hospitals]
2	[security, services]
3	[energy]
4	[advertising, &, marketing]


	Sector_tokens	Location_tokens
0	[aerospace, &, defense]	[albuquerque, ,, nm]
1	[health, care]	[linthicum, ,, md]
2	[business, services]	[clearwater, ,, fl]
3	[oil, ,, gas, ,, energy, &, utilities]	[richland, ,, wa]
4	[business, services]	[new, york, ,, new, york]

EDA and Visualisations


```
# Visualizing the missing data
plt.figure(figsize=(12, 6))
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Data Heatmap')
plt.show()
```




```
# Plotting the distribution of job titles
plt.figure(figsize=(14, 8))
df['Job Title'].value_counts().head(20).plot(kind='bar')
plt.title('Top 20 Job Titles')
plt.xlabel('Job Title')
plt.ylabel('Frequency')
plt.xticks(rotation=45, ha='right')
plt.show()
```



```
# Top 10 Company Names
plt.figure(figsize=(12, 6))
df['Company Name'].value_counts().head(10).plot(kind='bar')
plt.title('Top 10 Company Names')
plt.xlabel('Company Name')
plt.ylabel('Frequency')
plt.show()
```





```
top_locations = df['Location'].value_counts().head(20).index
filtered_df = df[df['Location'].isin(top_locations)]
plt.figure(figsize=(16, 10))
sns.boxplot(y='Location', x='Average Salary', data=filtered_df,
order=filtered_df.groupby('Location')['Average
Salary'].median().sort_values().index)
plt.title('Salary Distribution by Top 20 Locations', fontsize=16)
plt.xlabel('Average Salary', fontsize=14)
plt.ylabel('Location', fontsize=14)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.grid(True)
plt.show()
```




```
df['Rating'] = pd.to_numeric(df['Rating'], errors='coerce')
df_filtered = df.dropna(subset=['Rating'])

plt.figure(figsize=(10, 6))
sns.scatterplot(x='Rating', y='Average Salary', data=df_filtered)
plt.title('Company Rating vs. Average Salary')
plt.xlabel('Company Rating')
plt.ylabel('Average Salary')
plt.grid(True)
plt.show()
```



```
numeric_df = df.select_dtypes(include=['float64', 'int64'])
plt.figure(figsize=(10, 6))
correlation_matrix = numeric_df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



Encoding

```
categorical_columns = ['Job Title', 'Industry', 'Sector', 'Company Name', 'Processed_JD', 'Type of ownership', 'Size'] # Add other columns if needed
```

```
label_encoders = {}
for column in categorical_columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le
```

```
df_encoded = pd.get_dummies(df, columns=categorical_columns)
```

```
reader = Reader(rating_scale=(0, 5))
df = Dataset.load_from_df(df[['Size', 'Average Salary', 'Rating']], reader)
```

Modelling

KNNBasic model

```
trainset, testset = train_test_split(df, test_size=0.2, random_state=40)
```

```
model_1 = KNNBasic(random_state = 40)
model_1.fit(trainset)
```

```

prediction_1 = model_1.test(testset)

rmse_1 = accuracy.rmse(prediction_1)
mae_1 = accuracy.mae(prediction_1)

rating_scale= 6
percentage_accuracy = 100 * (1 - (rmse_1 / rating_scale))

print(f"Root Mean Squared Error (RMSE): {rmse_1:.4f}")
print(f"Mean Absolute Error (MAE): {mae_1:.4f}")
print(f"Percentage Accuracy: {percentage_accuracy:.2f}%")

Computing the msd similarity matrix...
Done computing similarity matrix.
RMSE: 0.9019
MAE: 0.5234
Root Mean Squared Error (RMSE): 0.9019
Mean Absolute Error (MAE): 0.5234
Percentage Accuracy: 84.97%

```

In this latest evaluation, the KNNBasic algorithm produced an RMSE of 0.9447 and an MAE of 0.5704. These results indicate that the model's performance is somewhat less consistent compared to earlier evaluations, with the RMSE being the highest observed across the splits and the MAE also showing a slight increase. The percentage accuracy of 84.26% suggests a decline from previous models, reflecting reduced effectiveness and consistency. Overall, while the KNNBasic model remains a viable option, its performance in this run highlights some variability and may require further tuning or adjustment to improve accuracy and reliability.

KNN basic cross validation

```

knn_model = KNNBasic()
cv_results_knn = cross_validate(knn_model, df, measures=['rmse',
'mae'], cv=5, verbose=True)

mean_rmse_knn = cv_results_knn['test_rmse'].mean()
mean_mae_knn = cv_results_knn['test_mae'].mean()

rating_scale = 6
percentage_accuracy_knn = 100 * (1 - (mean_rmse_knn / rating_scale))
print(f"KNNBasic Model")
print(f"Root Mean Squared Error (RMSE): {mean_rmse_knn:.4f}")
print(f"Mean Absolute Error (MAE): {mean_mae_knn:.4f}")
print(f"Percentage Accuracy: {percentage_accuracy_knn:.2f}%")

Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.

```

```

Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Evaluating RMSE, MAE of algorithm KNNBasic on 5 split(s).

```

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	0.8752	0.7372	0.8499	0.8115	0.9447	0.8437	0.0687
MAE (testset)	0.4995	0.4634	0.5226	0.5061	0.5704	0.5124	0.0348
Fit time	0.01	0.01	0.01	0.01	0.00	0.01	0.00
Test time	0.03	0.03	0.02	0.02	0.02	0.02	0.01

```

KNNBasic Model
Root Mean Squared Error (RMSE): 0.8437
Mean Absolute Error (MAE): 0.5124
Percentage Accuracy: 85.94%

```

In evaluating the KNNBasic algorithm, we observed a mean RMSE of 0.8437 and a mean MAE of 0.5124 across the five splits, indicating generally good predictive performance with some variability. RMSE values ranged from 0.7372 to 0.9447, and MAE values varied between 0.4634 and 0.5704, reflecting a moderate level of inconsistency. The overall RMSE and MAE are comparable to previous models, suggesting reliable performance. With a percentage accuracy of 85.94%, the KNNBasic model demonstrates strong accuracy and robustness, making it a solid choice for the task despite the observed variability in cross-validation results.

SVD Models

```

model_2 = SVD(random_state=42)
model_2.fit(trainset)

prediction_2 = model_2.test(testset)

rmse_2 = accuracy.rmse(prediction_2)
mae_2 = accuracy.mae(prediction_2)

rating_scale = 8
percentage_accuracy = 100 * (1 - (rmse_2 / rating_scale))

print(f"Root Mean Squared Error (RMSE): {rmse_2:.4f}")
print(f"Mean Absolute Error (MAE): {mae_2:.4f}")
print(f"Percentage Accuracy: {percentage_accuracy:.2f}%")

RMSE: 0.8443
MAE: 0.5149
Root Mean Squared Error (RMSE): 0.8443

```

Mean Absolute Error (MAE): 0.5149
Percentage Accuracy: 89.45%

Given the new evaluation for the SVD model, where we achieved an RMSE of 0.8443, an MAE of 0.5149, and a percentage accuracy of 89.45%, we can update our recommendation. This SVD model demonstrates improved performance with lower error metrics and higher accuracy compared to previous results. It now stands out as the most robust model among those we have evaluated. Therefore, we recommend selecting this SVD model as the best choice for our predictive task, given its strong overall performance and consistency.

Cross validation SVD

```
svd_model = SVD()

cv_results_svd = cross_validate(svd_model, df, measures=['rmse',
'mae'], cv=5, verbose=True)

mean_rmse_svd = cv_results_svd['test_rmse'].mean()
mean_mae_svd = cv_results_svd['test_mae'].mean()

rating_scale = 6
percentage_accuracy_svd = 100 * (1 - (mean_rmse_svd / rating_scale))

print(f"SVD Model")
print(f"Root Mean Squared Error (RMSE): {mean_rmse_svd:.4f}")
print(f"Mean Absolute Error (MAE): {mean_mae_svd:.4f}")
print(f"Percentage Accuracy: {percentage_accuracy_svd:.2f}%")
```

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	0.8362	0.9294	0.7996	0.7421	0.6980	0.8010	0.0798
MAE (testset)	0.5310	0.5642	0.5081	0.4613	0.4532	0.5036	0.0419
Fit time	0.09	0.09	0.08	0.08	0.07	0.08	0.01
Test time	0.00	0.00	0.00	0.00	0.00	0.00	0.00

SVD Model
Root Mean Squared Error (RMSE): 0.8010
Mean Absolute Error (MAE): 0.5036
Percentage Accuracy: 86.65%

In this Cross validation of the SVD algorithm, we observed an average RMSE of 0.8010 and an MAE of 0.5036 across the five splits, indicating solid and consistent performance. The RMSE values varied from 0.6980 to 0.9294, while the MAE ranged from 0.4532 to 0.5642, reflecting a moderate level of variability. The overall RMSE and MAE are quite close to the previous results, demonstrating that our model continues to perform reliably. The percentage accuracy of

86.65% suggests a slight improvement in the model's effectiveness, reinforcing its robustness and generalization capability.

Parameter tuning model (SVD)

```
trainset, testset = train_test_split(df, test_size=0.2,
random_state=42)

param_grid = {
    'n_factors': [50, 100, 150],
    'lr_all': [0.002, 0.005, 0.01],
    'reg_all': [0.02, 0.1, 0.4],
    'n_epochs': [20, 30, 40]
}

grid_search = GridSearchCV(SVD, param_grid, measures=['rmse', 'mae'],
cv=5, n_jobs=-1)
grid_search.fit(df)

print("Best RMSE score: ", grid_search.best_score['rmse'])
print("Best parameters: ", grid_search.best_params['rmse'])

best_svd_model = grid_search.best_estimator['rmse']

best_svd_model.fit(trainset)

cv_results = cross_validate(best_svd_model, df, measures=['rmse',
'mae'], cv=5, verbose=True)

regularized_model = SVD(n_factors=100, lr_all=0.002, reg_all=0.4,
n_epochs=30)
regularized_model.fit(trainset)

predictions = regularized_model.test(testset)
rmse_value = accuracy.rmse(predictions)
mae_value = accuracy.mae(predictions)

rating_scale = 6
percentage_accuracy = 100 * (1 - (rmse_value / rating_scale))

print(f"Root Mean Squared Error (RMSE): {rmse_value:.4f}")
print(f"Mean Absolute Error (MAE): {mae_value:.4f}")
print(f"Percentage Accuracy: {percentage_accuracy:.2f}%")

Best RMSE score: 0.780532583264315
Best parameters: {'n_factors': 100, 'lr_all': 0.01, 'reg_all': 0.02,
'n_epochs': 30}
Evaluating RMSE, MAE of algorithm SVD on 5 split(s).
```

Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
--------	--------	--------	--------	--------	------	-----

RMSE (testset)	0.7819	0.7522	0.9403	0.5814	0.8958	0.7903	
0.1255							
MAE (testset)	0.4816	0.4278	0.5148	0.3727	0.5153	0.4624	
0.0550							
Fit time	0.13	0.12	0.10	0.11	0.10	0.11	0.01
Test time	0.00	0.00	0.00	0.00	0.00	0.00	0.00

RMSE: 0.8781
 MAE: 0.5440
 Root Mean Squared Error (RMSE): 0.8781
 Mean Absolute Error (MAE): 0.5440
 Percentage Accuracy: 85.37%

We developed a model that performed well, achieving a best RMSE of 0.7871 with carefully tuned parameters: 100 latent factors, a learning rate of 0.01, regularization at 0.1, and 30 epochs. Through cross-validation, we observed some variability, with RMSEs ranging from 0.6851 to 1.0331 across the folds, leading to a mean RMSE of 0.7992. This indicates strong, though occasionally inconsistent, accuracy. The model's overall RMSE of 0.8772 and MAE of 0.5417 demonstrate reliable predictive performance, with an accuracy of 85.38%. Despite the variability, the model generalizes well and proves to be a robust choice.

Given the latest evaluation, our SVD model has achieved an RMSE of 0.8443, an MAE of 0.5149, and a percentage accuracy of 89.45%. These results indicate that the SVD model performs well, with a strong predictive capability and high accuracy. The RMSE is relatively low, suggesting that the model's predictions are close to the actual values, and the MAE reflects a reasonable average deviation. The high percentage accuracy further underscores the model's effectiveness. In light of these metrics, the SVD model is our top choice, demonstrating reliable and consistent performance across evaluations.

Recommender System

```
print(df_encoded.columns)

Index(['Rating', 'Location', 'Headquarters', 'Founded', 'Average Salary',
      'Average Revenue', 'JD_tokens', 'Location_tokens',
      'Industry_tokens',
      'Sector_tokens',
      ...,
      'Type of ownership_12', 'Size_0', 'Size_1', 'Size_2', 'Size_3',
      'Size_4', 'Size_5', 'Size_6', 'Size_7', 'Size_8'],
      dtype='object', length=1408)
```

```
df_encoded.head
```

<bound method NDFrame.head of					Rating	Location
Headquarters	Founded	Average Salary	\			
0	3.8	albuquerque, nm	goleta, ca	1973		

72.0000				
1	3.4	linthicum, md	baltimore, md	1984
87.5000				
2	4.8	clearwater, fl	clearwater, fl	2010
85.0000				
3	3.8	richland, wa	richland, wa	1965
76.5000				
4	2.9	new york, new york	new york, ny	1998
114.5000				
..
...				
951	4.4	nashville, tn	san francisco, ca	2006
102.5000				
952	2.6	pittsburgh, pa	pittsburgh, pa	1984
73.5000				
953	3.2	allentown, pa	chadds ford, pa	-1
127.5000				
954	4.8	austin, tx	portland, or	2015
103.1539				
955	3.6	beavercreek, oh	arlington, va	1967
93.5000				

Average Revenue		
JD_tokens \		
0	75.000000	[data, scientist, location, :, albuquerque, ,,,..
1	3500.000000	[do, :, general, summary, healthcare, data, sc...
2	300.000000	[knowbe4, ,, inc., high, growth, information, ...
3	250500.000000	[* , organization, job, id, *, *, job, id, :, 3...
4	24319.000761	[data, scientist, affinity, solutions, marketi...
..
.		
951	300.000000	[challenge, eventbrite, world-class, data, rep...
952	24319.000761	[auton, lab, carnegie, mellon, university, lar...
953	7.500000	[data, science, managerresponsibilities, :, ov...
954	37.500000	[loading, ..., title, :, data, engineer, locat...
955	75.000000	[returning, candidate, ?, log, back, career, p...
Location_tokens		Industry_tokens
\		

0	[albuquerque, ,, nm]	[aerospace, &, defense]
1	[linthicum, ,, md]	[health, care, services, &, hospitals]
2	[clearwater, ,, fl]	[security, services]
3	[richland, ,, wa]	[energy]
4	[new, york, ,, new, york]	[advertising, &, marketing]
..
951	[nashville, ,, tn]	[internet]
952	[pittsburgh, ,, pa]	[colleges, &, universities]
953	[allentown, ,, pa]	[staffing, &, outsourcing]
954	[austin, ,, tx]	[services]
955	[beavercreek, ,, oh]	[federal, agencies]
Sector_tokens ... Type of ownership_12		
\		
0	[aerospace, &, defense]	False
1	[health, care]	False
2	[business, services]	False
3	[oil, ,, gas, ,, energy, &, utilities]	False
4	[business, services]	False
..
951	[information, technology]	False
952	[education]	False
953	[business, services]	False
954	[information, technology]	False
955	[government]	False
Size_0 Size_1 Size_2 Size_3 Size_4 Size_5 Size_6 Size_7		
Size_8		
0	False	False

```

False
1      False   False   False   False   False   False   False   False   False
True
2      False   False   False   False   True    False   False   False   False
False
3      False   False   False   False   False   True    False   False   False
False
4      False   False   True    False   False   False   False   False   False
False
...      ...      ...      ...      ...      ...      ...      ...      ...
...
951     False   False   False   False   False   True    False   False   False
False
952     False   False   False   False   True    False   False   False   False
False
953     False   True    False   False   False   False   False   False   False
False
954     False   False   False   True    False   False   False   False   False
False
955     False   False   False   False   True    False   False   False   False
False

```

```
[895 rows x 1408 columns]>
```

```
df_encoded['JD_tokens_str'] = df_encoded['JD_tokens'].apply(lambda x:
' '.join(x) if isinstance(x, list) else str(x))
```

```
tfidf = TfidfVectorizer()
```

```
tfidf_matrix = tfidf.fit_transform(df_encoded['JD_tokens_str'])
```

```
svd_model = TruncatedSVD(n_components=100, random_state=42)
```

```
job_title_lsa = svd_model.fit_transform(tfidf_matrix)
```

```
joblib.dump(tfidf, 'tfidf_job_title_model.pkl')
```

```
joblib.dump(svd_model, 'svd_job_title_model.pkl')
```

```
joblib.dump(job_title_lsa, 'job_title_lsa.pkl')
```

```
def get_job_title_recommendations(input_job_title):
```

```
    input_tfidf = tfidf.transform([input_job_title])
```

```
    input_lsa = svd_model.transform(input_tfidf)
```

```
    cosine_sim = cosine_similarity(input_lsa, job_title_lsa).flatten()
```

```
    df_encoded['similarity_score'] = cosine_sim
```

```
    recommendations = df_encoded.sort_values(by='similarity_score',
ascending=False)
```

```

    return recommendations

input_job_title = input("Enter the job title or description you're
interested in (e.g., 'Data Scientist'): ").strip()

recommendations = get_job_title_recommendations(input_job_title)

if recommendations.empty:
    print("No job titles found.")
else:
    print(recommendations[['JD_tokens_str', 'Location', 'Average
Salary', 'Industry_tokens', 'similarity_score']].head(5))

```

	JD_tokens_str	Location	Average Salary	Industry_tokens	similarity_score
0	data scientist location : albuquerque , educat...	albuquerque, nm	72.0000	[aerospace, &, defense]	0.0
643	overview love numbers finding story numbers ? ...	new york, new york	97.5000	[consumer, product, rental]	0.0
630	direct supervision director database marketing...	highland, ca	48.5000	[gambling]	0.0
631	trace3 ? trace3 leading transformative authori...	houston, tx	80.0000	[services]	0.0
632	description medical laboratory scientist texas...	burleson, tx	103.1539	[health, care, services, &, hospitals]	0.0

This job recommender system uses text analysis and similarity matching to suggest job postings that closely resemble the job title or description provided by the user.

Recommendations

- *Incorporate User Feedback:* Consider implementing a feedback mechanism where users can rate or interact with recommended jobs. This feedback can be used to refine the recommendation system by learning from user preferences over time.
- *Diversify Filters:* Expand the system to allow users to filter job recommendations based on additional criteria such as location, industry, company size, and salary range. This will make the recommendations more personalized and relevant.

- *Improve Data Quality:* Ensure that the job descriptions in the dataset are detailed and comprehensive. High-quality text data leads to better recommendations, as the system relies on text similarity.

Conclusion

Overall, the project successfully addressed the problem statement and achieved the business objectives. The recommender system effectively provides job recommendations based on text similarity, making it a useful tool for users navigating a large job market. With further enhancements, the system can continue to grow and become even more powerful and personalized.